

# dog\_app

January 10, 2021

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with **‘(IMPLEMENTATION)’** in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a **‘TODO’** statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **‘Question X’** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **‘Answer:’**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.


The rubric contains *optional* “Stand Out Suggestions” for enhancing the project beyond the minimum requirements. If you decide to pursue the “Stand Out Suggestions”, you should include the code in this Jupyter notebook.

#### 1.1.1 Why We’re Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied

image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).

```
hello, dog!  
your predicted breed is ...  
American Staffordshire terrier
```



Sample Dog Output

In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

### 1.1.2 The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- Section ??: Import Datasets
- Section ??: Detect Humans
- Section ??: Detect Dogs
- Section ??: Create a CNN to Classify Dog Breeds (from Scratch)
- Section ??: Create a CNN to Classify Dog Breeds (using Transfer Learning)
- Section ??: Write your Algorithm
- Section ??: Test Your Algorithm

---

#### ## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: \* Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location `/dogImages`.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location `/lfw`.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

## 1.2 Get data on Colab

I used Colab to run my project, the following cells are just to set up the environment.

## 1.3 Mount Google Drive

```
[ ]: from google.colab import drive
drive.mount('/content/drive')

%cd /content
!ls -alh
```

Drive already mounted at `/content/drive`; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

```
/content
total 24K
drwxr-xr-x 1 root root 4.0K Jan  9 18:17 .
drwxr-xr-x 1 root root 4.0K Jan  9 18:14 ..
drwxr-xr-x 1 root root 4.0K Jan  6 18:10 .config
drwx----- 5 root root 4.0K Jan  9 18:17 drive
drwxr-xr-x 7 root root 4.0K Jan  9 18:19 project-dog-classification
drwxr-xr-x 1 root root 4.0K Jan  6 18:10 sample_data
```

## 1.4 Get the data

```
[ ]: %mkdir project-dog-classification
%cd project-dog-classification

! cp -a /content/drive/MyDrive/Udacity/deep-learning-v2-pytorch/
  ↳project-dog-classification/* .

! wget https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/dogImages.zip
! wget https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/lfw.zip

! unzip -qq dogImages.zip
! rm dogImages.zip

! unzip -qq lfw.zip
! rm lfw.zip
! rm -rf __MACOSX

! ls -alh
```

```
mkdir: cannot create directory project-dog-classification: File exists
/content/project-dog-classification
cp: cannot open '/content/drive/MyDrive/Udacity/deep-
learning-v2-pytorch/project-dog-classification/Report.gdoc' for reading:
Operation not supported
```

## 1.5 Check GPU assigned

```
[ ]: import torch

# check if CUDA is available
use_cuda = torch.cuda.is_available()

if use_cuda:
    print ("Training on GPU...")
else:
    print ("Training on CPU...")

# check GPU type
gpu_info = !nvidia-smi
gpu_info = '\n'.join(gpu_info)
if gpu_info.find('failed') >= 0:
    print('Select the Runtime > "Change runtime type" menu to enable a GPU_
→accelerator, ')
    print('and then re-execute this cell.')
else:
    print(gpu_info)
```

Training on GPU...

Sun Jan 10 02:10:32 2021

```
+-----+
| NVIDIA-SMI 460.27.04      Driver Version: 418.67      CUDA Version: 10.1      |
+-----+-----+-----+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |     MIG M.     |
+-----+-----+-----+-----+-----+-----+
|   0   Tesla V100-SXM2...    Off   | 00000000:00:04:0 Off  |             0      |
| N/A   39C    P0      24W / 300W |  10MiB / 16130MiB |      0%      Default  |
|                                       |                    |     ERR!     |
+-----+-----+-----+-----+-----+-----+

+-----+
| Processes:
| GPU   GI    CI          PID    Type    Process name                        GPU Memory
|       ID    ID                                   |          Usage
+-----+-----+-----+-----+-----+-----+
|
```

```
| No running processes found |
+-----+
|
```

## 1.6 Check memory

```
[ ]: from psutil import virtual_memory
ram_gb = virtual_memory().total / 1e9
print('Your runtime has {:.1f} gigabytes of available RAM\n'.format(ram_gb))

if ram_gb < 20:
    print('To enable a high-RAM runtime, select the Runtime > "Change runtime_
→type"')
    print('menu, and then select High-RAM in the Runtime shape dropdown. Then, ')
    print('re-execute this cell.')
else:
    print('You are using a high-RAM runtime!')
```

Your runtime has 27.4 gigabytes of available RAM

You are using a high-RAM runtime!

## 1.7 Project resumes here

```
[ ]: import numpy as np
      from glob import glob

      # load filenames for human and dog images
      human_files = np.array(glob("lfw/*/"))
      dog_files = np.array(glob("dogImages/*/"))

      # print number of images in each dataset
      print('There are %d total human images.' % len(human_files))
      print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

## ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
[ ]: import cv2
import matplotlib.pyplot as plt
%matplotlib inline
```

```

# extract pre-trained face detector
face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.
→xml')

# load color (BGR) image
img = cv2.imread(human_files[0])
# convert BGR image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# find faces in image
faces = face_cascade.detectMultiScale(gray)

# print number of faces detected in the image
print('Number of faces detected:', len(faces))


# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.7.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
[ ]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

### 1.7.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:** 99% for humans, 14% for dogs

```
[ ]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

n_humans_detected = 0
n_dogs_detected_as_humans = 0
for human, dog in zip(human_files_short, dog_files_short):
    if face_detector(human):
```

```

        n_humans_detected += 1
    if face_detector(dog):
        n_dogs_detected_as_humans += 1

print('Percentage human faces detected: {:.2f}'.format(n_humans_detected / 100.
→))
print('Percentage dogs detected as human faces: {:.2f}'.
→format(n_dogs_detected_as_humans / 100.))

```

Percentage human faces detected: 0.99

Percentage dogs detected as human faces: 0.14

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```

[ ]: ### (Optional)
    ### TODO: Test performance of another face detection algorithm.
    ### Feel free to use as many code cells as needed.

```

---

## ## Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

### 1.7.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```

[ ]: import torch
    import torchvision.models as models

    # define VGG16 model
    VGG16 = models.vgg16(pretrained=True)

    # check if CUDA is available
    use_cuda = torch.cuda.is_available()

    if use_cuda:
        print ("Training on GPU...")
    else:
        print ("Training on CPU...")

    # move model to GPU if CUDA is available

```



```

if use_cuda:
    VGG16 = VGG16.cuda()

#print(VGG16)

```

Training on GPU...

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

### 1.7.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher\_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

[ ]: from PIL import Image
import torchvision.transforms as T

# Set PIL to be tolerant of image files that are truncated.
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """
    # pre-process images
    transform = T.Compose([T.RandomResizedCrop(224),
                           T.ToTensor(),
                           T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
→0.224, 0.225])])

    image = Image.open(img_path)
    x = transform(image)[np.newaxis, :]
    if use_cuda:
        x = x.cuda()
    return torch.argmax(VGG16(x))

```

### 1.7.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```
[ ]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):

    p = VGG16_predict(img_path)
    return p >= 151 and p <= 268 # ImageNet dog's categories as per_
    →documentation above
```

### 1.7.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

**Answer:**

Percentage humans detected as dog: 2% Percentage dogs detected as dogs: 97%

```
[ ]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.

n_humans = 0
n_dogs = 0

for human, dog in zip(human_files_short, dog_files_short):
    if dog_detector(human):
        n_humans += 1
    if dog_detector(dog):
        n_dogs += 1

print('Percentage humans detected as dog: {:.2f}'.format(n_humans / 100.))
print('Percentage dogs detected as dogs: {:.2f}'.format(n_dogs / 100.))
```

Percentage humans detected as dog: 0.02

Percentage dogs detected as dogs: 0.97

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
[ ]: ### (Optional)
### TODO: Report the performance of another pre-trained network.
```

### Feel free to use as many code cells as needed.

---

### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

#### 1.7.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
[ ]: import torch

from torchvision import datasets
from torchvision import transforms

def get_mean_std(transforms_array, image_dir):
    """
    return the mean and std for the images in the specified dir passed in,
    after applying the specified transforms
    """
    transform = T.Compose(transforms_array + [transforms.ToTensor()])
    dataset = datasets.ImageFolder(image_dir, transform=transform)

    means = []
    stds = []
    for img, _ in dataset:
        means.append(torch.mean(img))
        stds.append(torch.std(img))

    mean = torch.mean(torch.tensor(means)).item()
    std = torch.mean(torch.tensor(stds)).item()
    return mean, std
```

```
[ ]: import os
import torch

from torchvision import datasets
from torchvision import transforms

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes

data_dir = 'dogImages/'
train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')
test_dir = os.path.join(data_dir, 'test/')

batch_size = 32
num_workers = 0

base_transforms = [transforms.Resize(160), transforms.CenterCrop(128)]
data_aug_transforms = [transforms.RandomHorizontalFlip(), transforms.
    ↪RandomRotation(10)]

#mean, std = get_mean_std(base_transforms, train_dir)
#print('Training data mean {:.3f}, std {:.4f}'.format(mean, std))
```

```

mean, std = (0.442, 0.2307)

normalize_transforms = [transforms.ToTensor(),
                        transforms.Normalize(mean=(mean, mean, mean), std=(std,
→std, std))]

train_transforms = transforms.Compose(base_transforms +
                                     data_aug_transforms +
                                     normalize_transforms)

# No data augmentation for valid and test data
valid_transforms = transforms.Compose(base_transforms + normalize_transforms)

train_data = datasets.ImageFolder(train_dir, transform=train_transforms)
valid_data = datasets.ImageFolder(valid_dir, transform=valid_transforms)
test_data = datasets.ImageFolder(test_dir, transform=valid_transforms)

print('Num training images: ', len(train_data))
print('Num valid images: ', len(valid_data))
print('Num test images: ', len(test_data))

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
→num_workers=num_workers,
→shuffle=True)

valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size,
→num_workers=num_workers,
→shuffle=False)

test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
→num_workers=num_workers,
→shuffle=False)

loaders_scratch = {'train' : train_loader,
                  'test': test_loader,
                  'valid': valid_loader}

```

```

Num training images: 6680
Num valid images: 835
Num test images: 836

```

```

[:]: # Visualize some sample data
import matplotlib.pyplot as plt
%matplotlib inline

# obtain one batch of training images
dataiter = iter(train_loader)

```

```

images, labels = dataiter.next()
images = images.numpy() # convert images to numpy for display

def imshow(img):
    img = img * std + mean # unnormalize
    plt.imshow(np.transpose(img, (1, 2, 0))) # convert from Tensor image

# plot the images in the batch, along with the corresponding labels
fig = plt.figure(figsize=(25, 4))
for idx in np.arange(20):
    ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
    imshow(images[idx])
    ax.set_title(train_data.classes[labels[idx]][5:])

```



**Question 3:** Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:**

The code first resizes the images to 160, a value slightly larger than the crop size, and then center-crops to 128. I picked 128 in order to have a size that is not too large but at the same time keeps a reasonable resolution, and is a power of 2. I resized first in order to avoid losing too much information from the picture.

I've augmented the data set with random horizontal flips and rotations of up to 10 degrees.

I've also calculated the mean and std of the training data in order to have more accurate normalization, rather than using the VGG parameters, which were obtained from a subset of ImageNet, or a more generic 0.5.

### 1.7.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

[ ]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()

```

```

        # Conv Layer block 1 - width/height after MP: 128 -> 64, channels: 3 ->
→32
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3,
→padding=1)

        # Conv Layer block 2 - width/height after MP: 64 -> 32, channels: 32 ->
→64
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3,
→padding=1)

        # Conv Layer block 3 - width/height after MP: 32 -> 16, channels: 64 ->
→128
        self.conv3 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3,
→padding=1)

        # Conv Layer block 4 - width/height after MP: 16 -> 8, channels: 128 ->
→256
        self.conv4 = nn.Conv2d(in_channels=128, out_channels=256,
→kernel_size=3, padding=1)

        self.maxPool = nn.MaxPool2d(kernel_size=2, stride=2)

        self.fc1 = nn.Linear(16384, 4096)
        self.fc2 = nn.Linear(4096, 4096)
        self.fc3 = nn.Linear(4096, 133)

        self.batchNorm = nn.BatchNorm1d(16384)
        self.dropOut = nn.Dropout(p=0.5)

        # use xavier weights initialization as recommended in forums
        #for m in self.modules():
        #    if isinstance(m, nn.Linear) or isinstance(m, nn.Conv2d):
        #        nn.init.xavier_uniform_(m.weight)

def forward(self, x):
    # conv layers
    x = self.maxPool(F.relu(self.conv1(x)))
    x = self.maxPool(F.relu(self.conv2(x)))
    x = self.maxPool(F.relu(self.conv3(x)))
    x = self.maxPool(F.relu(self.conv4(x)))

    #print(x.shape)

    # flatten + batch norm
    x = x.view(x.size(0), -1)
    x = self.batchNorm(x)

```

```

        x = self.dropOut(x)
        x = F.relu(self.fc1(x))

        x = self.dropOut(x)
        x = F.relu(self.fc2(x))

        x = self.dropOut(x)
        x = self.fc3(x)
        #print(x.shape)

    return x

##-## You do NOT have to modify the code below this line. ##-##

# instantiate the CNN
model_scratch = Net()
print(model_scratch)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()
    print(torch.cuda.memory_summary(device=None, abbreviated=False))

```

```

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv4): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (maxPool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (fc1): Linear(in_features=16384, out_features=4096, bias=True)
  (fc2): Linear(in_features=4096, out_features=4096, bias=True)
  (fc3): Linear(in_features=4096, out_features=133, bias=True)
  (batchNorm): BatchNorm1d(16384, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (dropOut): Dropout(p=0.5, inplace=False)
)
=====|
|               PyTorch CUDA memory summary, device ID 0               |
|-----|
|          CUDA OOMs: 0          |          cudaMalloc retries: 0          | | | |
|---|---|---|---|---|
| Metric          | Cur Usage | Peak Usage | Tot Alloc | Tot Freed |
|-----|
| Allocated memory |    853 MB |    853 MB | 47403 MB | 46549 MB |
| from large pool |    852 MB |    852 MB | 46290 MB | 45438 MB |

```



from small pool	1 MB	5 MB	1112 MB	1111 MB
-----				
Active memory	853 MB	853 MB	47403 MB	46549 MB
from large pool	852 MB	852 MB	46290 MB	45438 MB
from small pool	1 MB	5 MB	1112 MB	1111 MB
-----				
GPU reserved memory	972 MB	972 MB	972 MB	0 B
from large pool	966 MB	966 MB	966 MB	0 B
from small pool	6 MB	6 MB	6 MB	0 B
-----				
Non-releasable memory	14580 KB	29123 KB	21440 MB	21425 MB
from large pool	14256 KB	26688 KB	19502 MB	19488 MB
from small pool	324 KB	3027 KB	1937 MB	1937 MB
-----				
Allocations	51	65	9649	9598
from large pool	16	28	5016	5000
from small pool	35	38	4633	4598
-----				
Active allocs	51	65	9649	9598
from large pool	16	28	5016	5000
from small pool	35	38	4633	4598
-----				
GPU reserved segments	17	17	17	0
from large pool	14	14	14	0
from small pool	3	3	3	0
-----				
Non-releasable allocs	3	9	4205	4202
from large pool	2	6	2604	2602
from small pool	1	4	1601	1600
=====				

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:**

I've created a architecture with 4 convolutional layers and 3 fully connected linear layers.

The convolutional layers use kernel size set to 3 and padding set to 1 so as to preserve the dimensions, whereas the max pool layers use kernel size and stride set to 2 so as to half the dimensions. The four convolutional layers therefore reduce the dimentisons from 128 to 8. The channels are instead increased from 3 to 256 with steps of 32, 64, 128 and 256.

The fully connected layers have an intermediate layer at 4096. I've also added batch normalization and drop out.

I experimented with initial parameters initialization but had no improvement and hence reverted to default initialization.

The dropout is quite high because I noticed overfitting, although that was before adding batch normalization so it's possible that it could be reduced somewhat.

I got 18% accuracy with this architecture and therefore stopped experimenting futher.

### 1.7.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
[ ]: import torch.optim as optim

### TODO: select loss function
criterion_scratch = nn.CrossEntropyLoss()

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr = 0.001)
#optimizer_scratch = optim.Adadelta(model_scratch.parameters())
```

### 1.7.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
[ ]: # the following import is required for training to be robust to truncated
    → images

from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

from datetime import datetime

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()

            ## find the loss and update the model parameters accordingly
            optimizer.zero_grad()
            output = model(data)
```

```

        loss = criterion(output, target)
        loss.backward()
        optimizer.step()

        ## record the average training loss
        ##train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))
        train_loss += loss.item() * data.size(0)

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    output = model(data)
    loss = criterion(output, target)
    #valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))
    valid_loss += loss.item()*data.size(0)

# calculate average losses
train_loss = train_loss/len(train_loader.sampler)
valid_loss = valid_loss/len(valid_loader.sampler)

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.
→format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    #print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model .
→...'.format(
        # valid_loss_min, valid_loss))
        #
        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss

# return trained model
return model

```

```
[ ]: # train the model
model_scratch = train(100, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

%cp model_scratch.pt /content/drive/MyDrive/Udacity/deep-learning-v2-pytorch/
→project-dog-classification

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))
```

Epoch: 1	Training Loss: 4.213442	Validation Loss: 4.275726
Epoch: 2	Training Loss: 4.209490	Validation Loss: 4.264143
Epoch: 3	Training Loss: 4.186931	Validation Loss: 4.241871
Epoch: 4	Training Loss: 4.166542	Validation Loss: 4.227767
Epoch: 5	Training Loss: 4.152678	Validation Loss: 4.213223
Epoch: 6	Training Loss: 4.122827	Validation Loss: 4.193310
Epoch: 7	Training Loss: 4.114234	Validation Loss: 4.175648
Epoch: 8	Training Loss: 4.089961	Validation Loss: 4.167672
Epoch: 9	Training Loss: 4.082011	Validation Loss: 4.147907
Epoch: 10	Training Loss: 4.043008	Validation Loss: 4.151161
Epoch: 11	Training Loss: 4.033876	Validation Loss: 4.120167
Epoch: 12	Training Loss: 4.015846	Validation Loss: 4.103840
Epoch: 13	Training Loss: 3.998983	Validation Loss: 4.094554
Epoch: 14	Training Loss: 3.980329	Validation Loss: 4.073943
Epoch: 15	Training Loss: 3.967782	Validation Loss: 4.078675
Epoch: 16	Training Loss: 3.948219	Validation Loss: 4.057493
Epoch: 17	Training Loss: 3.937939	Validation Loss: 4.052180
Epoch: 18	Training Loss: 3.904809	Validation Loss: 4.024696
Epoch: 19	Training Loss: 3.899450	Validation Loss: 4.011857
Epoch: 20	Training Loss: 3.892359	Validation Loss: 4.038985
Epoch: 21	Training Loss: 3.867366	Validation Loss: 4.020607
Epoch: 22	Training Loss: 3.841542	Validation Loss: 3.989343
Epoch: 23	Training Loss: 3.833877	Validation Loss: 3.957730
Epoch: 24	Training Loss: 3.809459	Validation Loss: 3.953715
Epoch: 25	Training Loss: 3.794811	Validation Loss: 3.930320
Epoch: 26	Training Loss: 3.770808	Validation Loss: 3.921543
Epoch: 27	Training Loss: 3.762427	Validation Loss: 4.013195
Epoch: 28	Training Loss: 3.751422	Validation Loss: 3.895144
Epoch: 29	Training Loss: 3.723111	Validation Loss: 3.920633
Epoch: 30	Training Loss: 3.717900	Validation Loss: 3.882482
Epoch: 31	Training Loss: 3.698409	Validation Loss: 3.905858
Epoch: 32	Training Loss: 3.674988	Validation Loss: 3.927136
Epoch: 33	Training Loss: 3.663818	Validation Loss: 3.856540
Epoch: 34	Training Loss: 3.665636	Validation Loss: 3.842077
Epoch: 35	Training Loss: 3.631317	Validation Loss: 3.828759
Epoch: 36	Training Loss: 3.625916	Validation Loss: 3.824209
Epoch: 37	Training Loss: 3.594484	Validation Loss: 3.880293

Epoch: 38	Training Loss: 3.575470	Validation Loss: 3.832881
Epoch: 39	Training Loss: 3.573345	Validation Loss: 3.838829
Epoch: 40	Training Loss: 3.571329	Validation Loss: 3.903978
Epoch: 41	Training Loss: 3.532394	Validation Loss: 3.778379
Epoch: 42	Training Loss: 3.537697	Validation Loss: 3.773724
Epoch: 43	Training Loss: 3.519698	Validation Loss: 3.765071
Epoch: 44	Training Loss: 3.511332	Validation Loss: 3.764916
Epoch: 45	Training Loss: 3.474725	Validation Loss: 3.730516
Epoch: 46	Training Loss: 3.462928	Validation Loss: 3.754481
Epoch: 47	Training Loss: 3.462978	Validation Loss: 3.715966
Epoch: 48	Training Loss: 3.428128	Validation Loss: 3.729748
Epoch: 49	Training Loss: 3.419787	Validation Loss: 3.707428
Epoch: 50	Training Loss: 3.397734	Validation Loss: 3.792447
Epoch: 51	Training Loss: 3.387007	Validation Loss: 3.657484
Epoch: 52	Training Loss: 3.369203	Validation Loss: 3.675496
Epoch: 53	Training Loss: 3.369098	Validation Loss: 3.659058
Epoch: 54	Training Loss: 3.338314	Validation Loss: 3.651053
Epoch: 55	Training Loss: 3.344469	Validation Loss: 3.641096
Epoch: 56	Training Loss: 3.310745	Validation Loss: 3.610890
Epoch: 57	Training Loss: 3.285555	Validation Loss: 3.630297
Epoch: 58	Training Loss: 3.282943	Validation Loss: 3.605442
Epoch: 59	Training Loss: 3.260860	Validation Loss: 3.589093
Epoch: 60	Training Loss: 3.248093	Validation Loss: 3.604250
Epoch: 61	Training Loss: 3.239595	Validation Loss: 3.556570
Epoch: 62	Training Loss: 3.214210	Validation Loss: 3.551086
Epoch: 63	Training Loss: 3.197039	Validation Loss: 3.569344
Epoch: 64	Training Loss: 3.175313	Validation Loss: 3.535691
Epoch: 65	Training Loss: 3.179964	Validation Loss: 3.644741
Epoch: 66	Training Loss: 3.151308	Validation Loss: 3.597156
Epoch: 67	Training Loss: 3.142454	Validation Loss: 3.725059
Epoch: 68	Training Loss: 3.128034	Validation Loss: 3.539426
Epoch: 69	Training Loss: 3.117319	Validation Loss: 3.581037
Epoch: 70	Training Loss: 3.108791	Validation Loss: 3.562429
Epoch: 71	Training Loss: 3.076000	Validation Loss: 3.501242
Epoch: 72	Training Loss: 3.078617	Validation Loss: 3.558385
Epoch: 73	Training Loss: 3.059233	Validation Loss: 3.488347
Epoch: 74	Training Loss: 3.054356	Validation Loss: 3.681603
Epoch: 75	Training Loss: 3.039527	Validation Loss: 3.543621
Epoch: 76	Training Loss: 3.022353	Validation Loss: 3.485885
Epoch: 77	Training Loss: 3.006807	Validation Loss: 3.496915
Epoch: 78	Training Loss: 3.000648	Validation Loss: 3.494311
Epoch: 79	Training Loss: 2.978202	Validation Loss: 3.469681
Epoch: 80	Training Loss: 2.952380	Validation Loss: 3.450315
Epoch: 81	Training Loss: 2.954918	Validation Loss: 3.440937
Epoch: 82	Training Loss: 2.939680	Validation Loss: 3.598203
Epoch: 83	Training Loss: 2.929989	Validation Loss: 3.455385
Epoch: 84	Training Loss: 2.911360	Validation Loss: 3.481536
Epoch: 85	Training Loss: 2.893280	Validation Loss: 3.471070

Epoch: 86	Training Loss: 2.883705	Validation Loss: 3.380924
Epoch: 87	Training Loss: 2.869690	Validation Loss: 3.451198
Epoch: 88	Training Loss: 2.855110	Validation Loss: 3.387254
Epoch: 89	Training Loss: 2.846319	Validation Loss: 3.395308
Epoch: 90	Training Loss: 2.823235	Validation Loss: 3.575538
Epoch: 91	Training Loss: 2.818140	Validation Loss: 3.479967
Epoch: 92	Training Loss: 2.803805	Validation Loss: 3.381556
Epoch: 93	Training Loss: 2.805637	Validation Loss: 3.440576
Epoch: 94	Training Loss: 2.772539	Validation Loss: 3.467289
Epoch: 95	Training Loss: 2.779428	Validation Loss: 3.387725
Epoch: 96	Training Loss: 2.758164	Validation Loss: 3.382800
Epoch: 97	Training Loss: 2.736392	Validation Loss: 3.429862
Epoch: 98	Training Loss: 2.728565	Validation Loss: 3.383054
Epoch: 99	Training Loss: 2.719252	Validation Loss: 3.417264
Epoch: 100	Training Loss: 2.678961	Validation Loss: 3.399731

```
[ ]: <All keys matched successfully>
```

### 1.7.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
[ ]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    with torch.no_grad():
        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['test']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            # forward pass: compute predicted outputs by passing inputs to the
            ↪model
            output = model(data)
            # calculate the loss
            loss = criterion(output, target)
            # update average test loss
            # test_loss += ((1 / (batch_idx + 1)) * (loss.data - test_loss))
            test_loss += loss.item()*data.size(0)
            # convert output probabilities to predicted class
            pred = output.data.max(1, keepdim=True)[1]
            # compare predictions to true label
```

```

        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).
→cpu().numpy())
        total += data.size(0)

    test_loss = test_loss/len(loaders['test'].sampler)
    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

```

```

[:]: # load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

# call test function
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.367971

Test Accuracy: 21% (177/836)

#### ## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

### 1.7.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

[:]: ## TODO: Specify data loaders for transfer learning solution, similar to what's
→done in step3, just change
## image preprocessing and increase batch size as per paper (https://arxiv.org/
→pdf/1409.1556.pdf)

import os
import torch

from torchvision import datasets
from torchvision import transforms

data_dir = 'dogImages/'
train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')

```

```

test_dir = os.path.join(data_dir, 'test/')

batch_size = 64
num_workers = 0

transform = T.Compose([T.Resize(256),
                        T.CenterCrop(224),
                        T.ToTensor(),
                        T.Normalize(mean=[0.485, 0.456, 0.406],
                                   std=[0.229, 0.224, 0.225])])

train_data = datasets.ImageFolder(train_dir, transform=transform)
valid_data = datasets.ImageFolder(valid_dir, transform=transform)
test_data = datasets.ImageFolder(test_dir, transform=transform)

print('Num training images: ', len(train_data))
print('Num valid images: ', len(valid_data))
print('Num test images: ', len(test_data))

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           num_workers=num_workers,
                                           →shuffle=True)

valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size,
                                           num_workers=num_workers,
                                           →shuffle=False)

test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers,
                                           →shuffle=False)

loaders_transfer = {'train': train_loader,
                   'test': test_loader,
                   'valid': valid_loader}

```

```

Num training images: 6680
Num valid images: 835
Num test images: 836

```

### 1.7.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

[:]: import torchvision.models as models
import torch.nn as nn

model_transfer = models.vgg19(pretrained=True)

```



```

print(model_transfer)

for param in model_transfer.features.parameters():
    param.requires_grad = False

n_inputs = model_transfer.classifier[6].in_features
last_layer = nn.Linear(n_inputs, 133)
model_transfer.classifier[6] = last_layer

print("Modified classifier: {}".format(model_transfer.classifier))

if use_cuda:
    model_transfer = model_transfer.cuda()

```

```

VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace=True)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace=True)
    (16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (17): ReLU(inplace=True)
    (18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
    (19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace=True)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (22): ReLU(inplace=True)
    (23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (24): ReLU(inplace=True)
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (26): ReLU(inplace=True)
    (27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,

```

```

ceil_mode=False)
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (29): ReLU(inplace=True)
    (30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (31): ReLU(inplace=True)
    (32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (33): ReLU(inplace=True)
    (34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (35): ReLU(inplace=True)
    (36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)
Modified classifier: Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=133, bias=True)
)

```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:**

Very deep neural networks with small filter should perform well with this task given that they were the winners for ImageNet in 2014, so I picked the deepest, VGG19 and decided to retrain the fully connected layers to specialize it on dog breeds.

#### 1.7.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```

[ ]: criterion_transfer = nn.CrossEntropyLoss()
optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.001)

```

### 1.7.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model\_transfer.pt'.

```
[ ]: # train the model

# The model parameters will be saved to this file
model_transfer = train(50, loaders_transfer, model_transfer,
    →optimizer_transfer, criterion_transfer, use_cuda, 'model_transfer.pt')
%cp model_transfer.pt /content/drive/MyDrive/Udacity/deep-learning-v2-pytorch/
    →project-dog-classification

# load the model that got the best validation accuracy
model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

Epoch: 1	Training Loss: 4.608812	Validation Loss: 4.015163
Epoch: 2	Training Loss: 3.688275	Validation Loss: 3.118357
Epoch: 3	Training Loss: 2.822104	Validation Loss: 2.268434
Epoch: 4	Training Loss: 2.100472	Validation Loss: 1.619323
Epoch: 5	Training Loss: 1.575391	Validation Loss: 1.213565
Epoch: 6	Training Loss: 1.234933	Validation Loss: 0.961211
Epoch: 7	Training Loss: 1.038029	Validation Loss: 0.809496
Epoch: 8	Training Loss: 0.900965	Validation Loss: 0.709751
Epoch: 9	Training Loss: 0.818150	Validation Loss: 0.641758
Epoch: 10	Training Loss: 0.738214	Validation Loss: 0.595038
Epoch: 11	Training Loss: 0.687683	Validation Loss: 0.553164
Epoch: 12	Training Loss: 0.639502	Validation Loss: 0.524457
Epoch: 13	Training Loss: 0.613993	Validation Loss: 0.499590
Epoch: 14	Training Loss: 0.573933	Validation Loss: 0.485077
Epoch: 15	Training Loss: 0.545640	Validation Loss: 0.466856
Epoch: 16	Training Loss: 0.516461	Validation Loss: 0.454195
Epoch: 17	Training Loss: 0.504500	Validation Loss: 0.437418
Epoch: 18	Training Loss: 0.483047	Validation Loss: 0.430763
Epoch: 19	Training Loss: 0.461202	Validation Loss: 0.416344
Epoch: 20	Training Loss: 0.444239	Validation Loss: 0.407965
Epoch: 21	Training Loss: 0.432330	Validation Loss: 0.404256
Epoch: 22	Training Loss: 0.426797	Validation Loss: 0.396348
Epoch: 23	Training Loss: 0.396784	Validation Loss: 0.392552
Epoch: 24	Training Loss: 0.401540	Validation Loss: 0.385789
Epoch: 25	Training Loss: 0.379725	Validation Loss: 0.383760
Epoch: 26	Training Loss: 0.373309	Validation Loss: 0.380584
Epoch: 27	Training Loss: 0.361192	Validation Loss: 0.373664
Epoch: 28	Training Loss: 0.362099	Validation Loss: 0.369391
Epoch: 29	Training Loss: 0.349253	Validation Loss: 0.369375
Epoch: 30	Training Loss: 0.332869	Validation Loss: 0.365187
Epoch: 31	Training Loss: 0.325682	Validation Loss: 0.358383
Epoch: 32	Training Loss: 0.323413	Validation Loss: 0.363829
Epoch: 33	Training Loss: 0.314641	Validation Loss: 0.360208

Epoch: 34	Training Loss: 0.305631	Validation Loss: 0.353269
Epoch: 35	Training Loss: 0.297706	Validation Loss: 0.353830
Epoch: 36	Training Loss: 0.286893	Validation Loss: 0.351141
Epoch: 37	Training Loss: 0.282620	Validation Loss: 0.350017
Epoch: 38	Training Loss: 0.281323	Validation Loss: 0.349990
Epoch: 39	Training Loss: 0.284985	Validation Loss: 0.342053
Epoch: 40	Training Loss: 0.274616	Validation Loss: 0.341172
Epoch: 41	Training Loss: 0.264713	Validation Loss: 0.341487
Epoch: 42	Training Loss: 0.262014	Validation Loss: 0.344237
Epoch: 43	Training Loss: 0.260041	Validation Loss: 0.341495
Epoch: 44	Training Loss: 0.248595	Validation Loss: 0.344406
Epoch: 45	Training Loss: 0.250119	Validation Loss: 0.340515
Epoch: 46	Training Loss: 0.241131	Validation Loss: 0.333723
Epoch: 47	Training Loss: 0.236796	Validation Loss: 0.330791

```

KeyboardInterrupt                                Traceback (most recent call
last)

<ipython-input-71-bebc742afbe5> in <module>()
      2
      3 # The model parameters will be saved to this file
----> 4 model_transfer = train(100, loaders_transfer, model_transfer,
optimizer_transfer, criterion_transfer, use_cuda, 'model_transfer.pt')
      5 get_ipython().magic('cp model_transfer.pt /content/drive/MyDrive/
Udacity/deep-learning-v2-pytorch/project-dog-classification')
      6

<ipython-input-56-702892e84b8d> in train(n_epochs, loaders, model,
optimizer, criterion, use_cuda, save_path)
      20     #####
      21     model.train()
----> 22     for batch_idx, (data, target) in enumerate(loaders['train']):
      23         # move to GPU
      24         if use_cuda:

/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py in
__next__(self)
      433         if self._sampler_iter is None:
      434             self._reset()
--> 435         data = self._next_data()
      436         self._num_yielded += 1
      437         if self._dataset_kind == _DatasetKind.Iterable and \

```

```

/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py in
↳ _next_data(self)
    473     def _next_data(self):
    474         index = self._next_index() # may raise StopIteration
--> 475         data = self._dataset_fetcher.fetch(index) # may raise
↳ StopIteration
    476         if self._pin_memory:
    477             data = _utils.pin_memory.pin_memory(data)

```

```

/usr/local/lib/python3.6/dist-packages/torch/utils/data/_utils/fetch.py
↳ in fetch(self, possibly_batched_index)
    42     def fetch(self, possibly_batched_index):
    43         if self.auto_collation:
---> 44             data = [self.dataset[idx] for idx in
↳ possibly_batched_index]
    45         else:
    46             data = self.dataset[possibly_batched_index]

```

```

/usr/local/lib/python3.6/dist-packages/torch/utils/data/_utils/fetch.py
↳ in <listcomp>(.0)
    42     def fetch(self, possibly_batched_index):
    43         if self.auto_collation:
---> 44             data = [self.dataset[idx] for idx in
↳ possibly_batched_index]
    45         else:
    46             data = self.dataset[possibly_batched_index]

```

```

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/folder.py in
↳ __getitem__(self, index)
    149         """
    150         path, target = self.samples[index]
--> 151         sample = self.loader(path)
    152         if self.transform is not None:
    153             sample = self.transform(sample)

```

```

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/folder.py in
↳ default_loader(path)
    186         return accimage_loader(path)
    187     else:
--> 188         return pil_loader(path)
    189

```

190

```
    /usr/local/lib/python3.6/dist-packages/torchvision/datasets/folder.py in
pil_loader(path)
    168         with open(path, 'rb') as f:
    169             img = Image.open(f)
--> 170             return img.convert('RGB')
    171
    172

    /usr/local/lib/python3.6/dist-packages/PIL/Image.py in convert(self,
mode, matrix, dither, palette, colors)
    871         """
    872
--> 873         self.load()
    874
    875         if not mode and self.mode == "P":

    /usr/local/lib/python3.6/dist-packages/PIL/ImageFile.py in load(self)
    249
    250             b = b + s
--> 251             n, err_code = decoder.decode(b)
    252             if n < 0:
    253                 break
```

KeyboardInterrupt:

### 1.7.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
[ ]: model_transfer.load_state_dict(torch.load('model_transfer.pt'))

test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.398667

Test Accuracy: 89% (745/836)

### 1.7.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
[ ]: ### TODO: Write a function that takes a path to an image as input  
### and returns the dog breed that is predicted by the model.  
  
from PIL import Image  
import torchvision.transforms as T  
  
# Set PIL to be tolerant of image files that are truncated.  
from PIL import ImageFile  
ImageFile.LOAD_TRUNCATED_IMAGES = True  
  
# list of class names by index, i.e. a name can be accessed like class_names[0]  
class_names = [item[4:].replace("_", " ") for item in train_data.classes]  
  
def predict_breed_transfer(img_path):  
    # load the image and return the predicted breed  
    image = Image.open(img_path)  
    x = transform(image)[np.newaxis, :]  
    if use_cuda:  
        x = x.cuda()  
    return class_names[torch.argmax(model_transfer(x))]
```

---

#### ## Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the face\_detector and dog\_detector functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

```

hello, human!

0
200
400
600
800
1000
1200
1400
0 500 1000

You look like a ...
Chinese_shar-pei

```



Sample Human Output

### 1.7.18 (IMPLEMENTATION) Write your Algorithm

```

[:]: ### TODO: Write your algorithm.
### Feel free to use as many code cells as needed.

def show_image(img_path):
    #print("Image path: {}".format(img_path))
    img = cv2.imread(img_path)
    cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    plt.imshow(cv_rgb)
    plt.show()

def run_app(img_path):
    if dog_detector(img_path):
        print("Hello k9!")
        show_image(img_path)
        print("Predicted dog breed is {}".
→format(predict_breed_transfer(img_path)))
    elif face_detector(img_path):
        print("Hello human!")
        show_image(img_path)
        print("Resembling dog breed is {}".
→format(predict_breed_transfer(img_path)))
    else:
        print("Hello who are you?")
        show_image(img_path)
        print("Resembling dog breed is {}".
→format(predict_breed_transfer(img_path)))

    print("\n")

```

---

## Step 6: Test Your Algorithm



In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

### 1.7.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

**Answer:** Object detection should be improved, for example by using some algorithms such as Mask or Faster R-CNN or even YOLO. One possible tutorial is here: [https://pytorch.org/tutorials/intermediate/torchvision\\_tutorial.html](https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html). At the moment, the object detection algorithm mistakenly classified a dog and a cat, 2 errors out of 6 is a bit too high.

The algorithm could print the top 5 breed probabilities, rather than only the top one. This would work much better for the two dogs tested, since they are both mixed-breed.

The learning rate and optimizer I used for training the networks, both the one from scratch and VGG19, could be tuned further. Dropout is possibly too high but I observed significant overfitting with a dropout of 0.1, so I chose 0.5 similarly to the VGG paper. I did not have any luck with explicit parameter initialization in the network from scratch.

```
[ ]: ## TODO: Execute your algorithm from Step 6 on
    ## at least 6 images on your computer.
    ## Feel free to use as many code cells as needed.

own_files = np.array(glob("own_images/*"))


## suggested code, below
for file in own_files:
    run_app(file)
```

Hello k9!




Predicted dog breed is Black and tan coonhound

Hello who are you?



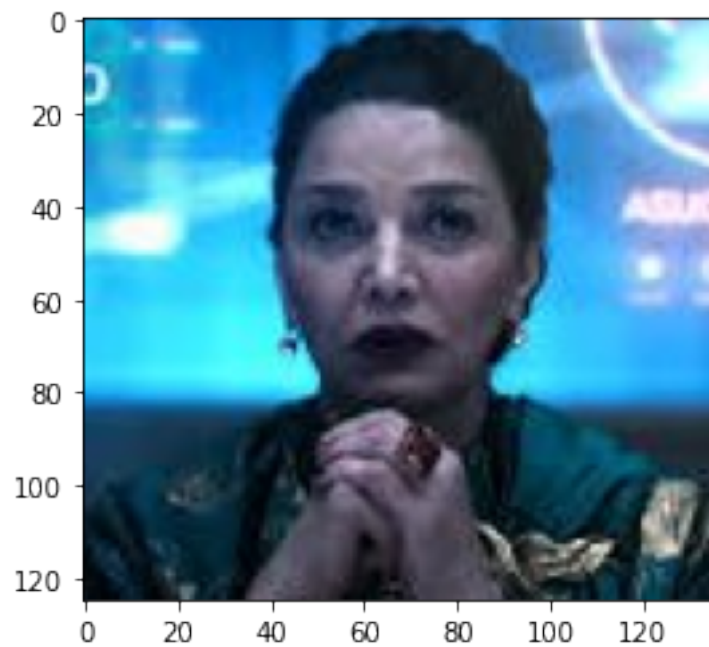
Resembling dog breed is Cardigan welsh corgi

Hello k9!




Predicted dog breed is Alaskan malamute

Hello human!



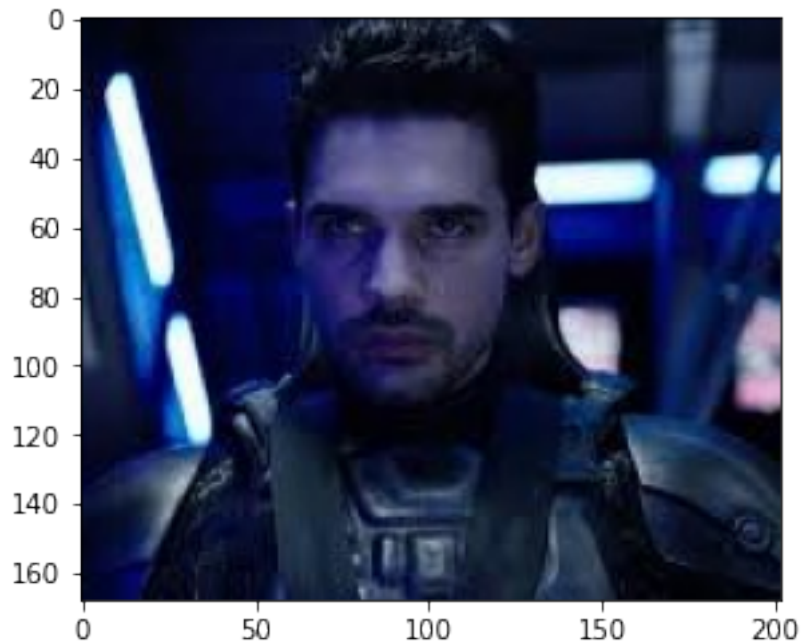
Resembling dog breed is Xoloitzcuintli

Hello who are you?




Resembling dog breed is Finnish spitz

Hello human!



Resembling dog breed is Basenji

Hello human!



Resembling dog breed is Collie